

## What is Wordpress

Wordpress is a content management system (CMS) and uses either MySQL or MariaDB database. I was able to create a website using Wordpress. Relevant AWS instances were used to create my Wordpress.

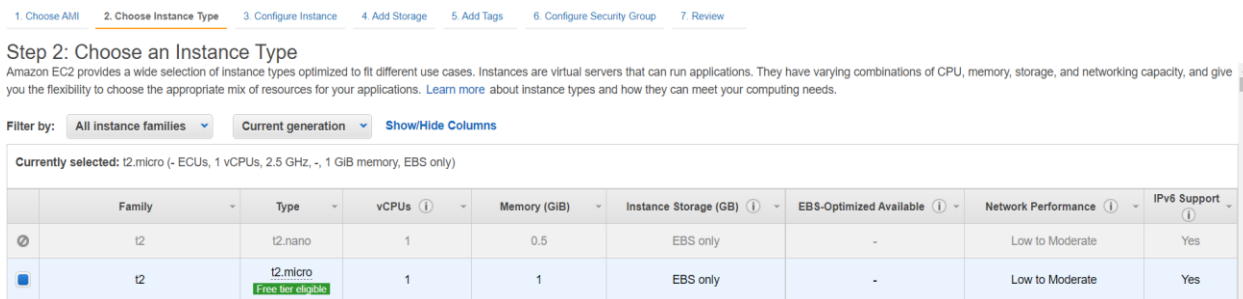
## Screenshots of my installation

### EC2


- 1) For AMI, I selected “Red Hat Enterprise Linux 8”



- 2) I created an EC2 with an Instance type of **t2.micro**.



- 3) For my Configure Instance Details, I created a subnet in us-east-1a.  
4) In the tag session, “**key: Name**” and “**Value: Jamie’s Website:**”  
5) For Security Group “**Type: HTTP**” and “**Source: Anywhere**”  
6) While creating my EC2, I created a new key pair to access my ec2 on putty. “**Key pair name: Jamie Key Pair**”



## Putty and Puttygen

- 1) Download putty and puttygen
- 2) Use puttygen to convert "JamieKeyPair.pem" to "JamieKeyPair.ppk"
- 3) After converting "JamieKeyPair.pem" to "JamieKeyPair.ppk", I can then now use this for my putty's private key.
- 4) For the "HostName", I paste the Public IPv4 address of my EC2 which is **75.101.155.191**
- 5) Hence I can then now use putty to download apache web server, httpd and create my wordpress database. The commands I used are :
  - a. Sudo yum install php-mysqld php-fpm mariadb-server httpd tar curl php-json
  - b. Sudo systemctl start mariadb
  - c. Sudo systemctl start httpd
  - d. Sudo systemctl enable mariadb
  - e. Sudo systemctl enable httpd
  - f. mysql\_secure\_installation
  - g. mysql -u root -p
  - h. CREATE DATABASE wordpress;
  - i. CREATE USER 'admin'@'localhost' IDENTIFIED BY 'pass';
  - j. GRANT ALL ON wordpress.\* TO 'admin'@'localhost';
  - k. FLUSH PRIVILEGES;
  - l. Exit
  - m. Sudo curl <https://wordpress.org/latest.tar.gz> --output wordpress.tar.gz
  - n. tar xf wordpress.tar.gz
  - o. sudo cp -r wordpress /var/www/html
  - p. sudo chown -R apache:apache /var/www/html/wordpress
  - q. sudo chcon -t httpd\_sys\_rw\_content\_t /var/www/html/wordpress -R

```
ec2-user@ip-172-31-31-97 ~]$ sudo dnf install php-mysqld php-fpm mariadb-server httpd tar curl php-json
Updating Subscription Management repositories.
Unable to read consumer identity

This system is not registered to Red Hat Subscription Management. You can use subscription-manager to register.

Last metadata expiration check: 1:39:32 ago on Thu 12 Aug 2021 05:21:08 AM UTC.
Package tar-2.1.30-5.el8.x86_64 is already installed.
Package curl-7.61.1-18.el8.x86_64 is already installed.
Dependencies resolved.
=====
Package Arch Version Repository Size
=====
Installing:
httpd x86_64 2.4.37-39.module+el8.4.0+9658+b87b2deb rhel-8-appstream-rhui-rpms 1.4 M
mariadb-server x86_64 3:10.3.28-1.module+el8.3.0+10472+7adc332a rhel-8-appstream-rhui-rpms 16 M
=====
```

```
ec2-user@ip-172-31-31-97:~
MariaDB [(none)]> CREATE DATABASE wordpress;
Query OK, 1 row affected (0.000 sec)

MariaDB [(none)]> CREATE USER 'admin'@'localhost' IDENTIFIED BY 'pass';
Query OK, 0 rows affected (0.000 sec)

MariaDB [(none)]> GRANT ALL ON wordpress.* TO 'admin'@'localhost';
Query OK, 0 rows affected (0.000 sec)

MariaDB [(none)]> FLUSH PRIVILEGES;
Query OK, 0 rows affected (0.000 sec)

MariaDB [(none)]> exit
Bye
ec2-user@ip-172-31-31-97 ~]$ sudo curl https://wordpress.org/latest.tar.gz --output wordpress.tar.gz
% Total % Received % Xferd Average Speed Time Time Time Current
Dload Upload Total Spent Left Speed
100 14.3M 100 14.3M 0 0 4357k 0 0:00:03 0:00:03 --:--:-- 4356k
ec2-user@ip-172-31-31-97 ~]$ sudo tar xf wordpress.tar.gz
ec2-user@ip-172-31-31-97 ~]$ sudo cp -r wordpress /var/www/html
ec2-user@ip-172-31-31-97 ~]$ sudo chown -R apache:apache /var/www/html/wordpress
ec2-user@ip-172-31-31-97 ~]$ sudo chcon -t httpd_sys_rw_content_t /var/www/html/wordpress -R
ec2-user@ip-172-31-31-97 ~]$
```

I can now access my wordpress website by using the public IPv4 address : <http://75.101.155.191/wordpress/>

# Wordpress

- 1) I set up the **Site Title, username, password and email** as follows:

## Welcome

Welcome to the famous five-minute WordPress installation process! Just fill in the information below and you'll be on your way to using the most extendable and powerful personal publishing platform in the world.

### Information needed

Please provide the following information. Don't worry, you can always change these settings later.

Site Title

Jamie's website

Username

Jamie

Names can have only alphanumeric characters, spaces, underscores, hyphens, periods, and the @ symbol.

Password

k5KRQEYB47BC#85Mmc

Hide

Strong

**Important:** You will need this password to log in. Please store it in a secure location.

Your Email

S10204749@connect.np.edu.sg

Double-check your email address before continuing.

Search engine visibility

☐ Discourage search engines from indexing this site

It is up to search engines to honor this request.

- 2) Hence, I can now then modify my website. As an example, I made my website look as if it was a portfolio

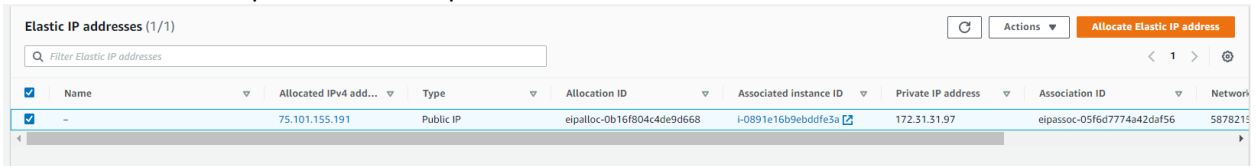


## Screenshots of enhancements and explanation why I added these enhancements

### Elastic IP address

I created this is so that when you restart the instance, it will restart with the same ip address. Hence, it is easier for me to go to my website as it would have the same public ipv4 address

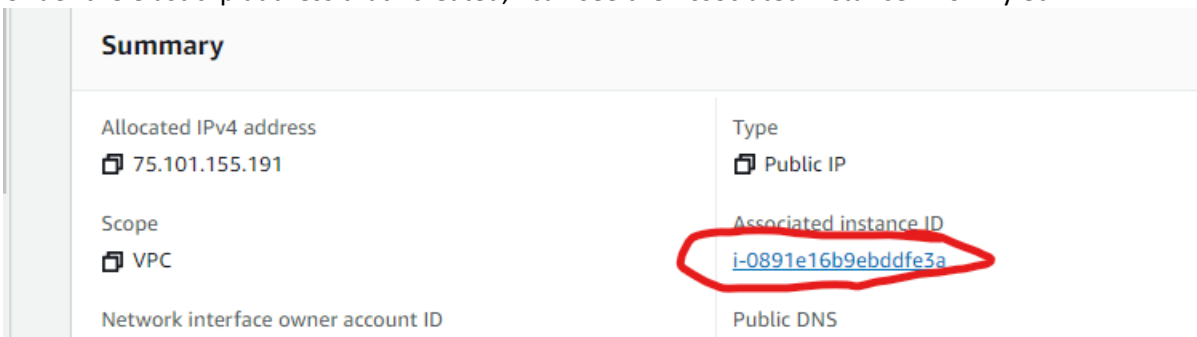
- 1) Click "Allocate Elastic IP address"
- 2) Public IPv4 address pool : Amazon's pool of IPv4 address



Elastic IP addresses (1/1)

<input checked="" type="checkbox"/>	Name	Allocated IPv4 address	Type	Allocation ID	Associated instance ID	Private IP address	Association ID	Network
<input checked="" type="checkbox"/>	-	75.101.155.191	Public IP	eipalloc-0b16f804c4de9d668	<a href="#">i-0891e16b9ebddfe3a</a>	172.31.31.97	eipassoc-05f6d7774a42daf56	5878215

- 3)
- 4) Under the elastic ip address that I created, I can see the Associated Instance ID of my ec2



Summary	
Allocated IPv4 address	Type
75.101.155.191	Public IP
Scope	Associated instance ID
VPC	<a href="#">i-0891e16b9ebddfe3a</a>
Network interface owner account ID	Public DNS

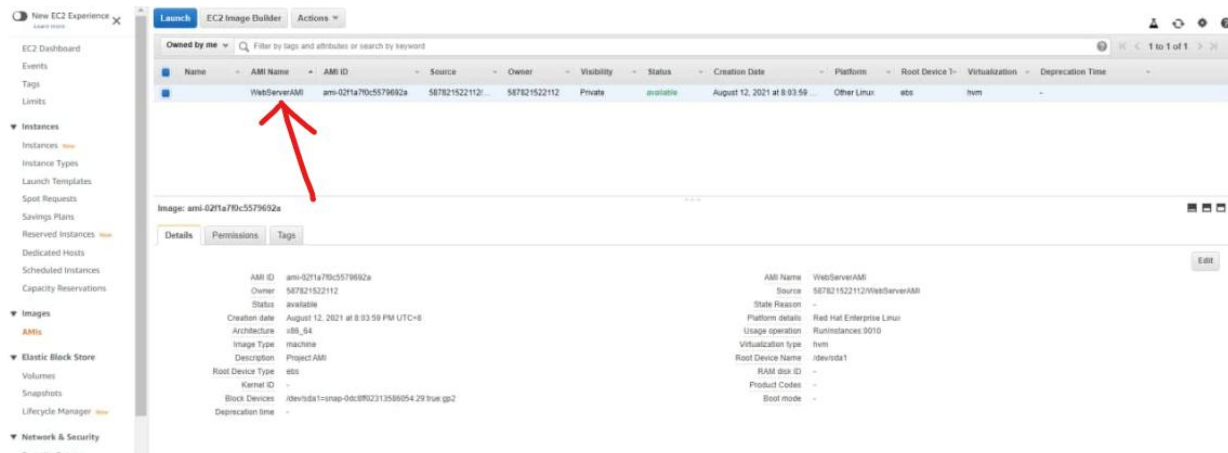


<input checked="" type="checkbox"/>	Jamie's website	<a href="#">i-0891e16b9ebddfe3a</a>	Running		t2.micro
-------------------------------------	-----------------	-------------------------------------	---------	--	----------

## AMI

This AMI is used when launching the Auto Scaling group. This is for me to have more instances with the same configuration.

- 1) Image name : WebServer AMI
- 2) Image description : Project AMI
- 3) Click "Create image"



The screenshot shows the AWS Management Console interface for the EC2 Image Builder service. On the left, there is a navigation menu with options like 'Instances', 'Images', 'Elastic Block Store', and 'Network & Security'. The main area displays a table of AMIs. A red arrow points to the 'WebServerAMI' entry. Below the table, the details for the selected image 'ami-021a70c5579692a' are shown, including its name, source, status, creation date, architecture, image type, description, root device type, kernel ID, block devices, and depreciation time.

Owned by me	Name	AMI Name	AMI ID	Source	Owner	Visibility	Status	Creation Date	Platform	Root Device	Virtualization	Deprecation Time
	WebServerAMI	ami-021a70c5579692a	ami-021a70c5579692a	587821522112	587821522112	Private	available	August 12, 2021 at 8:03:59	Other Linux	efs	hvm	-

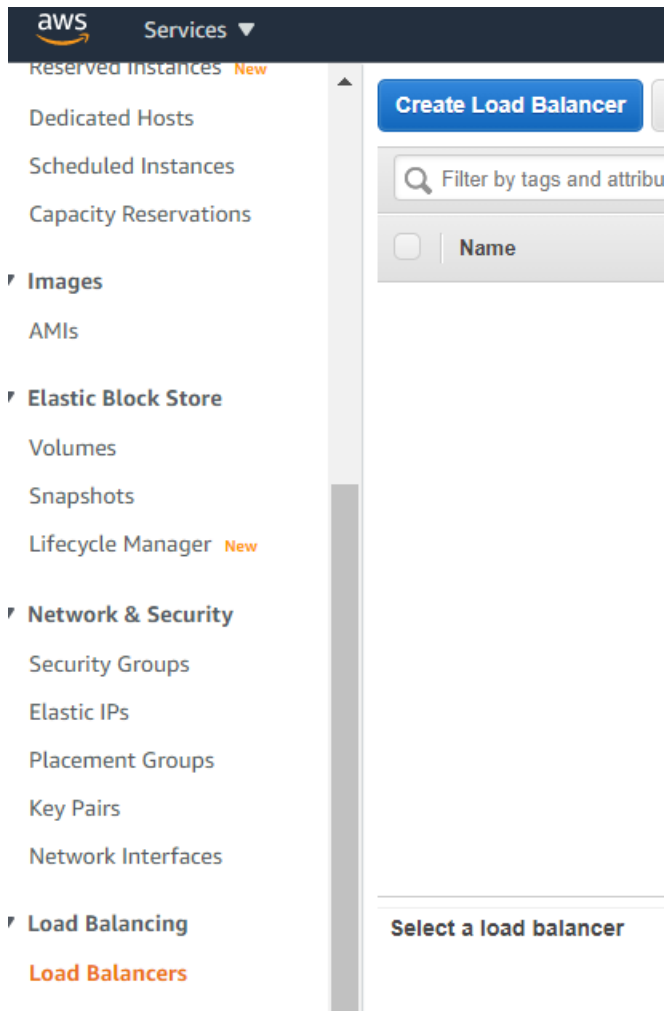
Image: ami-021a70c5579692a

Details	Permissions	Tags	
AMI ID	ami-021a70c5579692a	AMI Name	WebServerAMI
Owner	587821522112	Source	587821522112/WebServerAMI
Status	available	State Reason	-
Creation date	August 12, 2021 at 8:03:59 PM UTC+8	Platform details	Red Hat Enterprise Linux
Architecture	x86_64	Usage operation	RunInstances 9010
Image Type	machine	Virtualization type	hvm
Description	Project AMI	Root Device Name	/dev/sda1
Root Device Type	efs	RAM disk ID	-
Kernel ID	-	Product Codes	-
Block Devices	/dev/sda1--map-9dc8f8d21358605429 true gp2	Boot mode	-
Deprecation time	-		

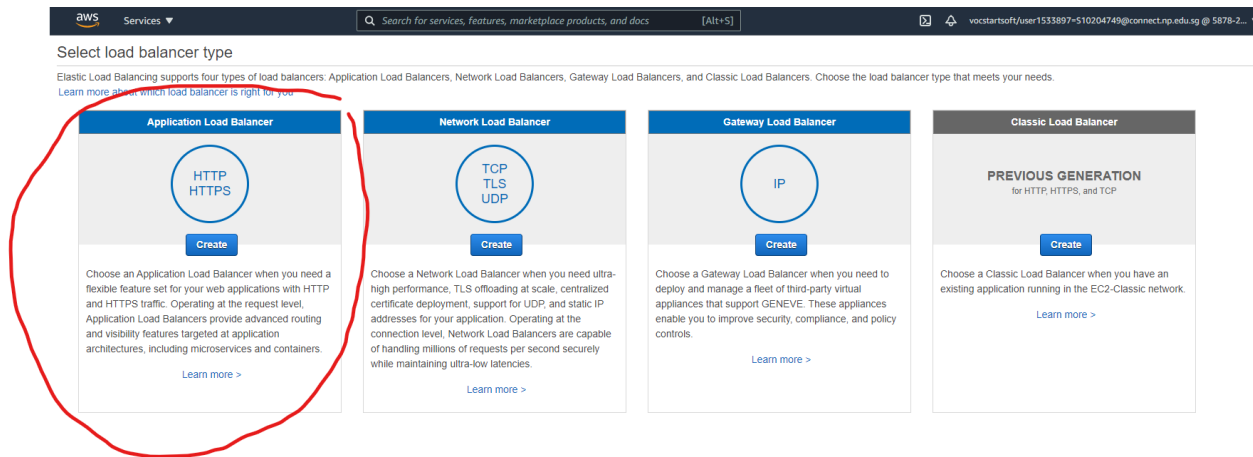
## Load balancer

When the load balancer identifies an unhealthy target, it redirects traffic to a healthy target. When it determines that the target is healthy again, it resumes directing traffic to it. Hence it enables be to have fault tolerance.

- 1) In the left navigation pane, click **Load balancer**.
- 2) Click **Create Load Balancer**



- 3) As the Load Balancer I chose **"HTTP HTTPS"**



- 4) I put the name as “WebServerELB”.
- 5) The load balancer protocol as HTTP.
- 6) For Availability Zone, I selected us-east-1a and us-east-1e.
- 7) I used the default VPC security group
- 8) I configured my routing as follows:

1. Configure Load Balancer

2. Configure Security Settings

3. Configure Security Groups

4. Configure Routing

5. Register Targets

6. Review

### Step 4: Configure Routing

Your load balancer routes requests to the targets in this target group using the protocol and port that you specify here. It also performs health checks on the targets or add listeners after the load balancer is created.

#### Target group

Target group ⓘ

New target group ↕

Name ⓘ

ProjectGroup

Target type

☒ Instance

☐ IP

☐ Lambda function

Protocol ⓘ

HTTP ↕

Port ⓘ

80

Protocol version ⓘ

☒ HTTP1

Send requests to targets using HTTP/1.1. Supported when the request protocol is HTTP/1.1 or HTTP/2.

☐ HTTP2

Send requests to targets using HTTP/2. Supported when the request protocol is HTTP/2 or gRPC, but gRPC-specific features are not available.

☐ gRPC

Send requests to targets using gRPC. Supported when the request protocol is gRPC.

#### Health checks

Protocol ⓘ

HTTP ↕

## Step 6: Review

Please review the load balancer details before continuing

### ▼ Load balancer

**Name** WebServerELB  
**Scheme** internet-facing  
**Listeners** Port:80 - Protocol:HTTP  
**IP address type** ipv4  
**VPC** vpc-e72c5c9a  
**Subnets** subnet-7761953b, subnet-bc3a518d  
**Tags**

### ▼ Security groups

**Security groups** sg-84e4eb9b

### ▼ Routing

**Target group** New target group  
**Target group name** ProjectGroup  
**Port** 80  
**Target type** instance  
**Protocol** HTTP  
**Protocol version** HTTP1  
**Health check protocol** HTTP  
**Path** /  
**Health check port** traffic port  
**Healthy threshold** 5  
**Unhealthy threshold** 2  
**Timeout** 5  
**Interval** 30

### Load Balancer Creation Status



#### Successfully created load balancer

Load balancer [WebServerELB](#) was successfully created.

Note: It might take a few minutes for your load balancer to be fully set up and ready to route traffic, and for the targets to complete the registration process and pass the initial health checks.

#### Suggested next steps

- Discover other services that you can integrate with your load balancer. Visit the **Integrated services** tab within [WebServerELB](#)
- Consider using AWS Global Accelerator to further improve the availability and performance of your applications. [AWS Global Accelerator console](#)

Close

Create Load Balancer

Actions

Filter by tags and attributes or search by keyword

1 to 1 of 1

Name	DNS name	State	VPC ID	Availability Zones	Type	Created At	Monitoring
WebServerELB	WebServerELB-1052801382...	Provisioning	vpc-e72c5c9a	us-east-1a, us-east-1e	application	August 12, 2021 at 8:23:58 ...	



## AUTO-SCALING GROUP

Amazon EC2 Auto Scaling adapts to changing conditions by adding or terminating instances, launching instances from an AMI, and ensuring that a minimum number of Amazon EC2 instances are operating.

- 1) In the left navigation pane, click **Launch Configurations**.
- 2) Click **“Create Launch configuration”**
- 3) I set up the **name, AMI and Instance type** as follows:

The screenshot shows the AWS Management Console interface for creating a launch configuration. The breadcrumb trail is 'EC2 > Launch configurations > Create launch configuration'. The main heading is 'Create launch configuration' with an 'Info' link. The form contains three sections: 'Launch configuration name' with a text input field containing 'Project Config'; 'Amazon machine image (AMI)' with a dropdown menu showing 'WebServerAMI'; and 'Instance type' with a dropdown menu showing 't2.micro (1 vCPUs, 1 GiB, EBS Only)' and a 'Choose instance type' button. Each of these three sections is circled in red.

- 4) I selected “Enable EC2 instance detailed monitoring within CloudWatch” for auto scaling to react quickly to changing utilization

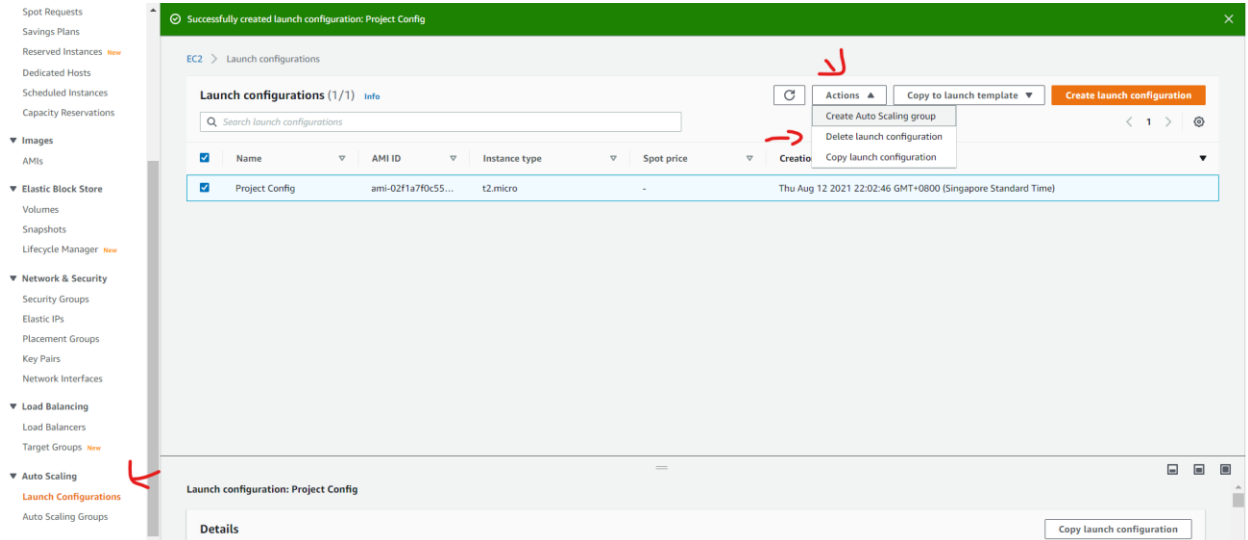
The screenshot shows the 'Additional configuration - optional' section of the AWS Management Console. It includes a 'Purchasing option' section with a 'Request Spot Instances' checkbox, an 'IAM instance profile' dropdown menu, and a 'Monitoring' section with a checked checkbox labeled 'Enable EC2 instance detailed monitoring within CloudWatch'. The 'Monitoring' section is circled in red. Below this is an 'EBS-optimized instance' section with a 'Launch as EBS-optimized instance' checkbox, and an 'Advanced details' section with a blue information box stating: 'Later, if you want to use a different launch configuration, you can create a new one and apply it to any Auto Scaling group. Existing launch configurations cannot be edited.'

5) Security group I used “launch-wizard-1”

<input checked="" type="checkbox"/>	Security group ID	Name	VPC ID	Description
<input type="checkbox"/>	sg-051880a52217725bb	launch-wizard-2	vpc-e72c5c9a	launch-wizard-2 created 2021-08-12T13:47:16.384+08:00
<input type="checkbox"/>	sg-0594d87ec47e84b9e	launch-wizard-3	vpc-e72c5c9a	launch-wizard-3 created 2021-08-12T14:16:37.086+08:00
<input checked="" type="checkbox"/>	sg-075d03bb64ea8aaa6	launch-wizard-1	vpc-e72c5c9a	launch-wizard-1 created 2021-08-12T12:27:28.701+08:00
<input type="checkbox"/>	sg-84e4eb9b	default	vpc-e72c5c9a	default VPC security group

6) I used an existing key pair which is “JamieKeyPair.pem”.


7) Under Launch configuration, I created an auto scaling group by clicking the following:



8) I set the name as “Project Auto Scaling Group”

9) I chose the subnets: “us-east-1a” and “us-east-1e”

10) I attached an existing load balancer which is the “Project Group|HTTP”



This will capture metrics at 1-minute intervals, which allows Auto Scaling to react quickly to changing usage patterns

11) I set the desired capacity, minimum capacity and maximum capacity as follows:

**Configure Settings**

Step 3 (optional)  
Configure advanced options

Step 4 (optional)  
**Configure group size and scaling policies**

Step 5 (optional)  
Add notifications

Step 6 (optional)  
Add tags

Step 7  
Review

### Group size - optional [Info](#)

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity  
1

Minimum capacity  
1

Maximum capacity  
2

### Scaling policies - optional

Choose whether to use a scaling policy to dynamically resize your Auto Scaling group to meet changes in demand. [Info](#)

☒ Target tracking scaling policy  
Choose a desired outcome and leave it to the scaling policy to add and remove capacity as needed to achieve that outcome.

☐ None

Scaling policy name  
ProjectScalingPolicy

Metric type  
Average CPU utilization

Target value  
60

Instances need  
300 seconds warm up before including in metric

This tells Auto Scaling to maintain an average CPU utilization across all instances 60%. Auto scaling will automatically add or remove capacity as required to keep the metric at, or close to, the specified target value. It adjusts to fluctuations in the metric due to a fluctuating load pattern

**Amazon CloudWatch**

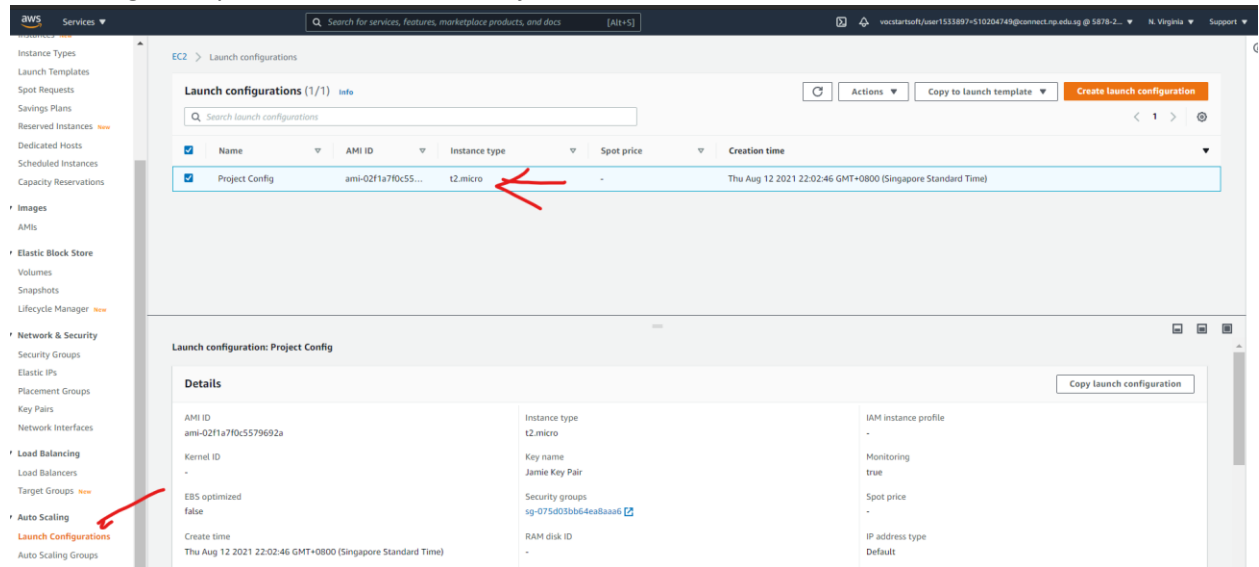
## CLOUD WATCH

This "OK" indicates that the CPU utilization has not exceeded the limit of 60% that I stated in autoscaling.

Name	State	Last state update	Conditions	Actions
TargetTracking-New Auto Scaling Group-AlarmLow-6362581b-8a1e-4b50-8d7a-c3c20ba76b80	In alarm	2021-09-12 23:17:40	CPUUtilization < 42 for 15 datapoints within 15 minutes	1 action
TargetTracking-New Auto Scaling Group-AlarmHigh-980ba295-4985-4e2c-a5d7-822a465c2995	OK	2021-09-12 23:05:14	CPUUtilization > 60 for 3 datapoints within 5 minutes	1 action

12)

### 13) For the tags : “Key: Name” and “Value: Project Instance”



## Difficulties I encountered

One of the difficulties I encountered is that I couldn't get to my EC2 instance using the command prompt. As an alternative, I decided to use PuTTY. While trying to connect with my PuTTY I realized that my key is in the wrong format: "JamieKeyPair.pem". Therefore, I downloaded PuTTYgen to convert the format from ".pem" to ".ppk" for it to be suitable to be used in PuTTY.

Furthermore, I didn't know how to edit my WordPress website. I was unsure of how to change the words and pictures that the WordPress automatically provides. After navigating WordPress and learning, I was able to create a nice portfolio in the end.

## Reflection

Throughout this module I learnt how to use the necessary services to create a website and use database. I gained resilience by finding new ways to solve the issues that I encountered. I realized that one of the important concepts I needed when working with AWS is that there would always be a way for my instances to not work properly. Hence, I had to learn how to use other services to tackle these problems. Using new services such as the Auto-scaling group, Load balancer and CloudWatch. After setting up the load balancer, I found a way to check if it was working by using its DNS name. By knowing if my services are working properly after creating, I became more confident in creating services in AWS. Furthermore, learning Linux commands have been a big help as I had to use the skills to download HTTP and Apache Web Server. While working through my way in AWS, I noticed that I also had to keep track of my spending as I was limited with \$50. Therefore, I became more careful when creating new instances and made sure that I stop my instances when I am not using them. To sum up, this project has helped me the most in understanding what is cloud computing and AWS in general.

